

Aegis: Automatic Enforcement of Security Policies in Workflow-driven Web Applications*

Luca Compagna
SAP Labs France
luca.compagna@sap.com

Serena Elisa Ponta
SAP Labs France
serena.ponta@sap.com

Daniel R. dos Santos
Fondazione Bruno Kessler
SAP Labs France
University of Trento
dossantos@fbk.eu

Silvio Ranise
Fondazione Bruno Kessler
ranise@fbk.eu

ABSTRACT

Organizations often expose business processes and services as web applications. Improper enforcement of security policies, e.g., authorization and control-flow integrity, in these applications leads to business logic vulnerabilities that are hard to find and may be devastating. AEGIS is a technique and tool to automatically synthesize run-time monitors to enforce control-flow and data-flow integrity, as well as authorization policies and authorization constraints in web applications. The enforcement of these properties can mitigate attacks such as forceful browsing, authorization bypass, and workflow violations, while allowing regulatory compliance in the form of, e.g., Separation of Duty. AEGIS also solves the satisfiability problem for constrained applications, i.e. determining if there is a possible execution of the application in which all tasks can be performed without violating the security policy. We evaluate AEGIS on a set of real-world benchmark applications, assessing the enforcement of policies, mitigation of vulnerabilities, and performance overhead.

Categories and Subject Descriptors

D.4.6 [Operating Systems]: Security and Protection; K.6.5 [Management of Computing and Information Systems]: Security and Protection

Keywords

Web Application; Policy Enforcement; Workflow Satisfiability

1. INTRODUCTION

Web applications are nowadays one of the preferred ways of exposing business processes and services to users. Many web

*This work has been partly supported by the EU under grant 317387 SECENTIS (FP7-PEOPLE-2012-ITN).

applications implement workflows, i.e. there is a pre-defined sequence of tasks that must be performed by users to reach a goal [3]. If an application does not correctly enforce its workflows, attackers can exploit this vulnerability to subvert it. In an e-commerce application, for instance, users must *Select products*, *Checkout*, *Enter shipping information*, *Enter payment information* and *Confirm*. If the application does not verify that user actions follow this sequence, a user can, e.g., skip the payment step and receive products without paying. Even errors generated while accessing pages in an unexpected order, can be exploited [27]. Workflow and business logic vulnerabilities are listed in the Common Weakness Enumeration (CWE)¹, in the OWASP Testing Guide [35] and (tangentially) in the OWASP Top 10 [38]. Control-flow integrity, i.e. the enforcement of an application's workflow, has been used in web applications to prevent workflow attacks and others, e.g., forceful browsing and race conditions [8, 9].

Data-flow integrity is also crucial and incorrect enforcement can lead to vulnerabilities where, e.g., a user can change the price of a product being purchased to pay less for it [53]. This kind of vulnerability is even more prominent in multi-party scenarios, where a user receives data such as tokens from one party (e.g., an identity provider) and must relay them to another party (e.g., a service provider). Several vulnerabilities have been discovered in recent years due to improper enforcement of data-flow integrity [53, 39, 44].

Besides control- and data-flow integrity, access control is fundamental for web application security whenever users must access only data and functionalities that they are authorized to by a given policy. Access control vulnerabilities are common and hard to find [45]. Moreover, some web applications implement collaborative work, in which many users work together to complete a workflow. Examples are Enterprise Resource Planning (ERP) software, used by employees in an organization to, e.g., manage purchases; and e-health applications, used by doctors and technical staff to manage patient records. In these applications, not only it is important to enforce authorization policies, but it may also be necessary to support authorization constraints, which impose more restrictions on what users can do at run-time. Examples of such constraints are Separation or Binding of Duty (BoD or SoD), requiring two different users (same user, respectively) to execute a pair of tasks; and cardinality, lim-

¹<https://goo.gl/yH4xrP> and <https://goo.gl/bFvzjK>

iting the number of tasks that a single user can execute in a workflow. These constraints can be used to avoid errors and frauds in security-critical applications that must follow regulatory compliance rules. Nonetheless, none of the applications we experimented with provided support for an easy to use, declarative specification of constraints. Including Odoo², an open-source ERP platform with more than 5000 developers and 2 million users, among them big companies such as Toyota and Danone. Without declarative specifications and proper enforcement, authorization constraints have to be implemented as application code embedded into each task [5] or translated to static assignments in the authorization policy. Both solutions are error-prone and can hardly scale.

Even with suitable specification and enforcement mechanisms, support for authorization policies and constraints may lead to situations where an application workflow cannot be completed because no user can execute an action without violating them. Determining if such a situation can be avoided, i.e. if a workflow can be completed in the presence of a policy and constraints, is known as the Workflow Satisfiability Problem (WSP) [49]. The WSP has received much attention in the workflow security community [29], but, to the best of our knowledge, has never been considered in web applications. In fact, transferring WSP solutions to the web domain is not trivial. These solutions often rely on a workflow model specification and a workflow management system to handle the control-flow of tasks and to provide an interface for users to request task executions, elements which are frequently not available for web applications.

In this paper, we present AEGIS³, a technique to synthesize run-time monitors for web applications that are capable of automatically (i) enforcing security policies composed of combinations of control- and data-flow integrity, authorization policies, and authorization constraints; and (ii) solving the run-time version of the WSP by granting or denying, at run-time, requests of users to perform tasks based on the satisfaction of the policy and constraints and the possibility to terminate the current workflow instance. AEGIS is based on [6], where a technique to synthesize run-time monitors that solve the WSP for security-sensitive workflow models was presented. We extend [6] by supporting data integrity. To synthesize a monitor, AEGIS first infers, using process mining [47], workflow models of the target application from a set of HTTP traces representing user actions. Traces must be manually edited to contain only actions that should be controlled by the monitor. Inferred models are Petri nets [36] labeled with HTTP requests representing tasks and annotated with data-flow properties obtained by using a set of heuristics based on differential analysis (as in, e.g., [53, 44]). These Petri nets (or a user-friendly representation, e.g., BPMN [52]) can be refined by a human user who, optionally, specifies authorization constraints and an authorization policy. A monitor is then generated from the model by pre-computing its possible executions [6]. At run-time, a reverse proxy is used to (i) capture login actions to later establish the acting users, and (ii) capture incoming requests and query the monitor to either allow or deny the request. AEGIS is completely black-box and can be used with new or legacy applications to support the enforcement of

security-related properties or to mitigate logic vulnerabilities.

The **main contribution** of this paper is the description and implementation of AEGIS, which paves the way for the support of complex security policies in workflow-driven web applications. **Another contribution** is an empirical evaluation on relevant applications, some taken from related work [39]. The rest of this paper is organized as follows: Section 2 presents an overview using motivating examples; Section 3 details the three steps of the technique; Section 4 shows the implementation and evaluation of our work; Section 5 discusses the approach and limitations; and Section 6 concludes the paper.

2. OVERVIEW

AEGIS synthesizes run-time monitors for workflow-driven web applications, i.e. applications implementing business processes and customer services as workflows. Hereafter, *web application* is used as an abbreviation for *workflow-driven web application*, unless stated otherwise.

A monitor synthesized by AEGIS can enforce three security-related properties: authorization policies (\mathcal{P}), defining which users are entitled to perform which tasks; authorization constraints (\mathcal{C}), defining run-time restrictions on the execution of tasks, e.g., a SoD requiring two different users to perform a pair of tasks; and control- and data-flow integrity (\mathcal{I}), specifying the authorized control-flow paths that the application must follow, as well as data invariants. Different web applications have different enforcement needs, which allows for the synthesis of different configurations of monitors, depending on which properties are switched on or off. We identify each configuration as a tuple containing the active properties, e.g., $\langle \mathcal{P}, \mathcal{C}, \mathcal{I} \rangle$, $\langle \mathcal{P}, \mathcal{I} \rangle$, $\langle \mathcal{C}, \mathcal{I} \rangle$, $\langle \mathcal{I} \rangle$. Control- and data-flow integrity are in the same category because it is not realistic that an application needs to enforce one and not the other.

AEGIS takes as input sets of HTTP traces representing user actions executed while interacting with a target web application. It synthesizes an external monitor composed of a set of queries to be used by a proxy sitting between users and the application. Each set of input traces is produced by a user simulating real clients completing a workflow as foreseen by the application (“good traces”). The monitor only enforces those workflows given as input by the user, having no impact on the rest of the application besides the overhead of a reverse proxy (which is frequently used in any case to implement, e.g., load balancing). Traces can be collected using test automation tools such as Selenium⁴ or ZAP⁵ and must be manually edited to contain only critical tasks. After trace collection, the whole technique is automated.

Figure 1 shows an overview of AEGIS. The top of the Figure shows the entire approach, where rectangles represent the three main steps (with sub-steps), yellow notes are inputs, and ovals are generated artifacts. The bottom of the Figure details the internals of the *Run-time Monitoring* component. The three main steps are the following.

1. Model Inference. The set of *HTTP traces* is automatically *stripped* of all information except request and response URLs, headers, and bodies; each request and response is *annotated* with *data-flow properties* inferred by a set of heuristics; traces are *aggregated* into a file called event log; and a process mining tool takes the log as input to generate

²<https://www.odoo.com/>

³Aegis was the mythological shield carried by Athena, and “under the aegis of” means “under the protection of.”

⁴<http://www.seleniumhq.org/>

⁵<http://goo.gl/XvxKd1>

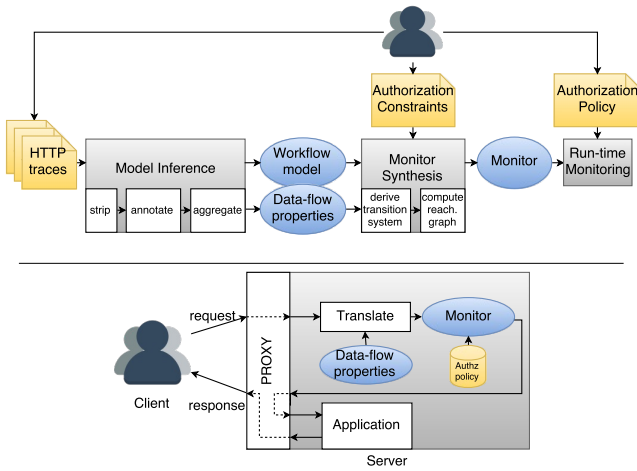


Figure 1: Overview of the technique

a Petri net *workflow model* whose transitions are labeled by the annotated requests. The inferred model can be refined according to the user’s understanding of the application.

2. Monitor Synthesis. Given a workflow model, the user specifies the *Authorization Constraints* to be enforced (if any) and whether an *Authorization Policy* will be provided at run-time. Control- and data-flow integrity are obtained automatically from the model, can be modified by the user, and are always enforced. The workflow model is presented in a convenient BPMN notation, and the specification of constraints is done graphically. A run-time monitor capable of enforcing the chosen properties is synthesized by translating the model to a symbolic *transition system* (the translation among BPMN, Petri nets, and transition systems is automatic [6]) and computing a *reachability graph*, whose nodes are labeled by first-order formulae describing states in the execution of the workflow. The graph represents all possible valid executions of the workflow by symbolic users, allowing us to support different authorization policies at run-time. The *Monitor* is a set of queries derived from the graph.

3. Run-time Monitoring. A reverse *proxy* is instantiated with the synthesized monitor and a concrete authorization policy (if required). It sits between users and the application, filters *requests* and *translates* them to the monitor. The monitor enforces the properties defined in step 2, granting a request if the control-flow is respected, the data-flow invariants hold, the user issuing the request is authorized by the policy, the authorization constraints are not violated and the current instance execution can still terminate. The proxy, based on the response from the monitor, may forward requests to the application or drop them to prevent the violation of some property.

A single application may implement several workflows. Steps 1 and 2 are performed for each workflow to be monitored, generating one monitor per workflow. Step 3 uses all the synthesized monitors and queries the correct one depending on the incoming request. Requests not related to any monitored workflow go directly to the application, without triggering a monitor query.

Below, we present two motivating examples that illustrate the configurations $\langle \mathcal{P}, \mathcal{C}, \mathcal{I} \rangle$, $\langle \mathcal{C}, \mathcal{I} \rangle$ (first example), and $\langle \mathcal{I} \rangle$ (second example). The first example motivates the distinctive

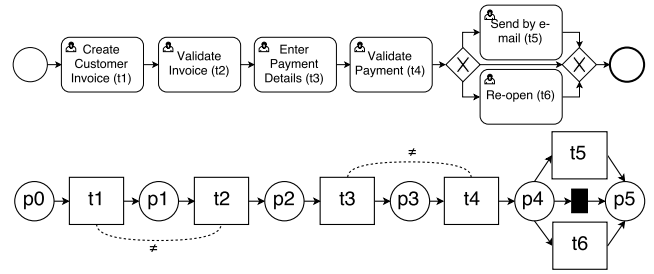


Figure 2: Customer invoice process in BPMN (top) and as a Petri net (bottom)

contributions of our technique: support for authorization policies, constraints, and workflow satisfiability. The second example shows that AEGIS is capable of mitigating logic vulnerabilities related to control- and data-flow integrity.

2.1 Example 1 - Enforcing constraints

Dolibarr⁶ is an open-source ERP web application for small and medium enterprises. It implements, among others, a business process similar to the one shown at the top of Figure 2 (in BPMN) to manage customer invoices.

The process contains 6 tasks (depicted by rounded boxes). Tasks $t1$ to $t4$ must be performed in sequence (as indicated by the solid arrows), while either $t5$, $t6$ or neither are performed last (as indicated by the diamond-shaped exclusive gateway). The original application implements each of the tasks shown in Figure 2. An authorization policy, control-flow, and possible data-flow invariants are implemented in an ad-hoc way, whose correctness is hard to verify, which may lead to vulnerabilities. The authorization policy originally supported by the application has a granularity of permissions that does not match the user-task assignment we support (there is no specific permission to, e.g., re-open an invoice or send it by e-mail). Authorization constraints are not supported. As a result, it is not trivial to prevent that a malicious user creates and validates a customer invoice (SoD between $t1$ and $t2$) or inserts and validates a payment (SoD between $t3$ and $t4$), which would allow him to, e.g., close invoices with an incorrect payment.

A user who wants to securely deploy this application can use AEGIS to generate a $\langle \mathcal{P}, \mathcal{C}, \mathcal{I} \rangle$ monitor to enforce control-flow integrity, ensuring that all the steps in the customer invoice process are performed in the correct order; an authorization policy, ensuring that only authorized users can execute each task; and the SoD constraints described above, to avoid frauds. If the user prefers to leave authorization enforcement to the application, a $\langle \mathcal{C}, \mathcal{I} \rangle$ monitor could be generated to only add support for constraints and integrity. To generate a monitor for the invoicing process, without impacting other parts of the application, the user starts by collecting traces simulating users performing the process. Some HTTP traces representing these executions are⁷:

```

 $\tau_1 = \{ /invoice?action=create&value=10&prod=abc,$ 
 $/invoice/validate?id=1, /invoice/pay/create?id=1&value=10,$ 
 $/invoice/pay/validate?id=1 \},$ 
 $\tau_2 = \{ /invoice?action=create&value=20&prod=def,$ 

```

⁶<http://www.dolibarr.org/>

⁷for the sake of readability, we show only simplified URLs, but headers and body are also part of the traces.

```

/invoice/validate?id=2,/invoice/pay/create?id=2&value=20,
/invoice/pay/validate?id=2,POST /invoice/send BODY id=2},
 $\tau_3 = \{ /invoice?action=create&value=30\&prod=ghi\&prod2=jk1,$ 
/invoice/validate?id=3,/invoice/pay/create?id=3&value=30,
/invoice/pay/validate?id=3,/invoice/reopen?id=3}.

```

Each trace τ_i represents one possible execution of the invoicing process and each request represents one task. The first four requests in each trace are essentially the same, but with different parameter values (e.g., `id` is 1 in τ_1 , 2 in τ_2 , and 3 in τ_3). They represent tasks t_1 , t_2 , t_3 , and t_4 . τ_1 is an example of the branch where only the first four tasks are executed, while t_5 is executed after t_4 in τ_2 , and t_6 is executed after t_4 in τ_3 . The traces are automatically analyzed to extract data-flow properties, annotated and aggregated into an event log, sent to a process mining tool and the resulting Petri net labeled with a task-to-URL map (**Step 1**). Figure 2 shows, at the bottom, the Petri net obtained from the process mining tool (ignore for a moment the dashed lines). The tasks in the net are labeled as t_i , with the following task-to-URL map:

```

 $t_1$  : /invoice?action=create&value=<<I>>&prod=<<DC>>,
 $t_2$  : /invoice/validate?id=<<IID>>,
 $t_3$  : /invoice/pay/create?id=<<IID>>&value=<<I>>,
 $t_4$  : /invoice/pay/validate?id=<<IID>>,
 $t_5$  : POST /invoice/send BODY id=<<IID>>,
 $t_6$  : /invoice/reopen?id=<<IID>>

```

Data-flow properties are represented by annotations on the URLs. The `<<IID>>` (instance identifier) annotation is applied to the elements used to bind all the requests to the same instance of a workflow, in this case the `id` parameter. The `<<I>>` (invariant) annotation is applied to values that should not change during the workflow, in this example the `value` of the invoice in t_1 should be the same as the `value` of the payment in t_2 . The `<<DC>>` (“don’t care”) annotation is applied to parameters that should be present in the request to help identify it as a unique action, but whose values are irrelevant. The parameter `prod2`, which is present in the request of t_1 only in τ_3 , is dropped in the task-to-URL map because it is considered optional, i.e. a trace may represent an invoice with one or more products, so only the first `prod` parameter needs to be present. These data-flow properties, as well as others not used in this example, are obtained by using heuristics detailed in Section 3.

The user then specifies the constraints that must be enforced, shown as dashed lines labeled by \neq in Figure 2. The model is used to synthesize a monitor (**Step 2**), which is composed of a set of SQL queries like

```

1 SELECT U2.ID FROM USERS AS U1, USERS AS U2, HST WHERE
2 HST.dt1 AND NOT HST.dt2 AND NOT HST.dt3 AND
3 NOT HST.dt4 AND NOT HST.dt5 AND NOT HST.dt6 AND
4 NOT HST.t1by = U1.ID AND U2.ID IN (SELECT * FROM
5 TA2) AND U1.ID IN (SELECT * FROM TA3) AND U2.ID IN
6 (SELECT * FROM TA4) AND U1.ID IN (SELECT * FROM TA5)

```

encoding that, to perform t_2 , only t_1 must have been executed (lines 2 and 3), there must be an authorized user u_2 who has not performed t_1 (line 4), and there must be other users capable of executing the remaining tasks (lines 5 and 6). This query is for the t_5 branch, there are similar queries for the t_6 branch and the branch where neither is executed, as well as other queries using a different number of users.

At run-time (**Step 3**), in the $\langle \mathcal{P}, \mathcal{C}, \mathcal{I} \rangle$ configuration, a policy is specified as a task-user assignment, e.g.,




Figure 3: Checkout process in BPMN (top) and as a Petri net (bottom)

$TA = \{(u_1, t_1), (u_1, t_2), (u_2, t_2), (u_3, t_3), (u_4, t_3), (u_4, t_4), (u_5, t_5), (u_6, t_6)\}$, where $(u, t) \in TA$ means that u is authorized to execute t . The assignment is stored in the database, and a reverse proxy is instantiated with the synthesized monitor. The proxy is capable of receiving a request such as

```
GET /invoice/validate?id=5 Cookie: sid=abcd1234
```

and identifying that it refers to task t_2 of instance 5 of the invoicing process being performed by user u_2 (whose cookie `sid`, sent in the header, has been stored during login). It then queries the monitor and, assuming that u_1 has previously executed t_1 and t_2 has not yet been executed, the query presented above is satisfied and the request is granted. On the other hand, a request can be blocked in several cases, such as if u_3 tries to execute t_2 ($(u_3, t_2) \notin TA$, i.e., not authorized by the policy), if u_1 tries to execute t_2 (there is a SoD between t_1 and t_2), if any user tries to execute t_3 before t_2 (because of the control-flow), or if any user issues a request for t_3 with a `value` different from the one sent for t_1 (because of the invariant).

To solve the WSP, regardless of the execution history, any request of u_4 to execute t_3 should also be blocked. Granting that request would mean that the only user authorized to execute t_4 has already executed t_3 , while both tasks are in SoD. Therefore, any execution where u_4 performs t_3 would either not terminate or terminate with the violation of some constraint or policy. This choice between business compliance (not violating policies) and business continuity (terminating the execution) should be avoided. The synthesized monitor presents a transparent way of avoiding it by blocking requests that lead to an undesired situation. Exceptional situations, where it is preferable to violate the policy with the knowledge of an administrator, can be accommodated by using a soft enforcement mode, as discussed in Section 5.

2.2 Example 2 - Mitigating vulnerabilities

TomatoCart⁸ is a popular e-commerce application that implements the checkout process depicted on the top of Figure 3. It is composed of 5 tasks executed in sequence, where t_4 is a sub-process that can be implemented in different ways, but must produce a data object representing a token issued by a trusted third party, that is read in t_5 .

This is an example of a multi-party web application [44], which implements the payment step by using third-party solutions such as PayPal⁹. An execution of this workflow, using PayPal Express Checkout, involves three actors: a client C , a service provider SP implementing TomatoCart and a trusted third party TTP implementing the payment provider. The execution starts with the client browsing the

⁸<http://www.tomatocart.com/>

⁹<https://www.paypal.com/>

SP, selecting some product ($t1$), requesting checkout ($t2$), and entering shipping information ($t3$). The *SP* then contacts the *TTP* and receives a token identifying the transaction (not shown in the workflow). The user is redirected to the *TTP* with the token ($t4$), completes the payment (again not shown in the Figure), and is redirected back to the *SP* passing the token, which is verified to complete the transaction ($t5$).

In version 1.1.7, TomatoCart had a vulnerability that allowed users to replay a PayPal Express Checkout token in $t5$ of a new transaction and shop for free [44]. This vulnerability was manually fixed in a later release of the application, but AEGIS could have been used to mitigate it until a patch was available (or until the patch could be applied, which is not always trivial). To mitigate the replay vulnerability, we can generate a monitor in the configuration $\langle \mathcal{I} \rangle$, enforcing control-flow integrity and the data invariant that the token received in $t4$ is the same that is sent in $t5$. An authorization policy and authorization constraints are not specified since every user can execute the steps in the checkout process and all steps are executed by the same user. Details of the communication between *SP* and *TTP* and between *C* and *TTP* are not shown in the workflow because the monitor only needs to enforce that no user can replace the token that has been sent to him/her. Although AEGIS ignores some messages, many vulnerabilities in multi-party web applications can be mitigated this way, because the vulnerabilities are commonly in the *SP* side [53].

To generate the monitor, we repeat the steps presented for Example 1. Below, there are some traces of the execution of the TomatoCart checkout process, again simplified for readability. Now the traces involve three parties, thus each request must be identified with its host.

```

 $\tau_1 = \{\text{shop.com/select, shop.com/checkout, shop.com/shipping,}$ 
 $\text{shop.com/payment} \rightarrow \text{paypal.com/webscr?token=abcd1234,}$ 
 $\text{shop.com/confirm?token=abcd1234}\},$ 
 $\tau_2 = \{\text{shop.com/select, shop.com/checkout, shop.com/shipping,}$ 
 $\text{shop.com/payment} \rightarrow \text{paypal.com/webscr?token=efgh5678,}$ 
 $\text{shop.com/confirm?token=efgh5678}\}.$ 

```

Figure 3 shows the Petri net obtained for the checkout process, labeled directly with the URL of each task (where \rightarrow represents a redirect). The invariant annotation $\langle \langle I \rangle \rangle$ is applied to the **token** received from PayPal, specifying that its value must be the same in **/payment** and **/confirm**.

A monitor is synthesized as before, however with neither authorization policy nor constraints. Workflow instances can be identified by the user identifier, since each user has only one checkout process at any given time. At run-time, whenever a user tries to replay a token, the monitor blocks this request because the token sent in $t5$ is different from the one received in $t4$ (since PayPal generates unique tokens). If the user tries to bypass the monitor by skipping step $t4$ and sending the token directly in $t5$, the monitor blocks the request because of a control-flow violation.

3. DETAILS

An HTTP trace (or a web session) is a sequence $S = \{(u_1 : r_1, s_1), (u_2 : r_2, s_2), \dots, (u_n : r_n, s_n)\}$ of pairs of web requests r_i issued by users u_i (which may or may not be all distinct) and responses s_i . Each web request or response is defined as $r_i = (\text{headers}, \text{body})$ and the information we derive from a request is a tuple $(\text{method}, \text{url}, P)$, where $\text{method} \in \{\text{GET}, \text{POST}\}$, url is the requested URL, and

P is a set of parameters of the form (k, v) , which can be in the URL (in GET requests), the body (in POST requests) or in the headers (e.g., cookies or *Location* in redirects). Data values passed as JSON can be flattened to the same representation. The parameters in P represent the data values later annotated with data-flow properties.

A workflow $W(T, U)$ is a set of tasks $(t \in T)$ in a causal order executed by a set of users $(u \in U)$, and a web application is composed of a set of workflows $\Psi = \{W_1(T_1, U_1), \dots, W_n(T_n, U_n)\}$. We take as input sets of web sessions $WS_i = \{S_1, S_2, \dots, S_n\}$ and infer from each WS_i a workflow $W_i(T_i, U_i)$, using a process mining function \mathcal{PM} , and a set of data property labels L_i , using heuristics. We also take as input, optionally, sets of authorization constraints C_i . We then use a monitor synthesis procedure $\mathcal{MS}(W_i, L_i, C_i)$ that returns a monitor M_i . M_i is capable of answering requests of the form “can user u perform task t ?”—encoded as $\text{can_do}(u, t)$ —with True iff the control-flow in W_i and the data-flow in L_i are respected, no authorization constraint in C_i is violated, the requesting user u is authorized by an authorization policy TA (specified at run-time), and the workflow can be executed until the end. TA is not taken as input by \mathcal{MS} because the procedure can accommodate different authorization policies given at run-time.

Attacks and enforcement. At run-time, a reverse proxy \mathcal{RP} receives an incoming request $u : r$ and based on the information taken from it, tries to translate it into a query of the form $\text{can_do}(u, t)$, for $u \in U_i$ and $t \in T_i$ of workflow $W_i(T_i, U_i)$, which can be answered by M_i . Attacks on the application at the level of web requests are reflected on the workflows [33] as shown below. The monitor can mitigate these attacks because they do not comply with the expected workflow (naturally, they are only mitigated in the parts of the application covered by the inferred model). A request forgery is an extra request not foreseen in a workflow $(\{r_1, r_2, \dots, r_{\text{forged}}, \dots, r_n\})$. A workflow bypass is a missing request $(\{r_1, r_2, \dots, r_{i-1}, r_{i+1}, \dots, r_n\})$. A workflow violation is an attempt to either repeat a unique request $(\{r_1, r_2, \dots, r_i, \dots, r_i, \dots, r_n\})$ or execute a request out of order $(\{r_1, r_2, \dots, r_{i+1}, \dots, r_i, \dots, r_n\})$. Authorization violations happen when a request is issued by a user who is not entitled to do so by the policy or when, for two tasks t_1 and t_2 in SoD, a user who previously issued a request r_1 to execute t_1 , issues a new request r_2 to execute t_2 .

Adversary model. The monitor enforces security properties related to access control and control- and data-flow integrity, ignoring vulnerabilities such as code injection. The target application is trusted, as well as any third parties trusted by it. Application users are not trusted, since they can be partially or fully controlled by an adversary.

3.1 Step 1 - Model inference

The goal of AEGIS is not to produce an accurate model of the whole application, but only workflow models containing a sequence of critical actions. Critical actions are the requests related to workflow tasks, whose execution should be controlled by the monitor. The definition of what is critical varies from application to application, but besides the usual noise in HTTP traces (e.g., requests loading images and other resources), any request that leaves the application state unchanged (e.g., AJAX requests for auto-completion of input fields) should be filtered out. Such requests are called navigation events, as opposed to system-interaction events,

which change the state of the application [42]. Not every system-interaction event should be controlled by the monitor (this should be decided by the user). However, discarding navigation events is crucial to keep the inferred models to a reasonable size and to eliminate imprecision due to variations in the process when executed by different users.

We assume that this treatment of the input traces is done before AEGIS is invoked. It can be done manually by deleting unimportant actions, but there are proposals of automated black-box techniques to detect state changing requests. Some techniques detect a state change by sending identical requests and comparing their outputs [17, 18], others are based on an abstraction of the user interface [42]. The former have limitations such as the need to isolate the application (other users cannot interact with it while a trace is collected) and to be able to reset it to its initial state. The latter cannot detect system states that are not reflected in the UI. Such techniques are usually embedded in crawlers to obtain a model of the entire application. Applying just the state-change detection part to traces of a single workflow may have sub-optimal results. Evaluating similar approaches to automate trace collection is left to future work.

Since some URLs in an application can take different parameters and different values for these parameters, while still representing the same action, and since we apply differential analysis to identify data-flow properties, we need at least two different traces as input, each containing a possible value for each of the parameters (including their presence and absence). The input traces should also represent all the possible execution paths of the process (control-flow). The number of input traces required for a precise model depends on the number of control-flow branches in the workflow being analyzed, as well as the diversity of the traces. Related works use, e.g., four traces as input [53] or traces with specific requirements for each of the parties in the process [44]. At least two login traces with distinct users must also be present, so that cookies defining the user session identifier and parameters representing user names can be mined, to map requests to concrete users at run-time.

It is possible to obtain input traces by reusing functional tests, which are common in web development and usually implemented using a framework such as Selenium. From the set of HTTP traces, we extract three artifacts: a workflow model, a task-to-URL map, and a set of data properties.

Workflow model and map. A workflow model is automatically obtained from a process mining tool. There are many well-known process mining algorithms and a simple example is the α -algorithm [47]. It mines workflow nets by recording all the events in a log and detecting relations between them, such as sequence (\rightarrow), exclusive or ($\#$), and parallel (\parallel). In the traces used in Example 1, it is possible to see that $t1$ always precedes $t2$ and $t2$ never precedes $t1$, so the algorithm infers a causal dependency between them and adds a place connecting transitions $t1$ and $t2$ in the output net (place $p1$ in Figure 2). It is also possible to see that $t4 \rightarrow t5$ ($t4$ precedes $t5$), $t4 \rightarrow t6$ ($t4$ precedes $t6$), and $t5 \# t6$ ($t5$ and $t6$ do not happen in the same trace), thus the algorithm creates a place after $t4$ that branches the execution ($p4$ in the same Figure). Since the input traces contain only relevant URLs and each unique URL becomes a transition after process mining, the task-to-URL map is trivial to obtain.

Data-flow properties. Identification and annotation of data properties has been used initially in [50] and later in

other works [53, 39, 44]. We use five annotations, namely *constant*, *don't care*, *invariant*, *instance identifier*, and *user identifier*, which are used for three goals. *Constants* and *don't cares* are used to restrict and generalize, respectively, the input traces by fixing or hiding given values that are used to match incoming requests at run-time. A *user identifier* is used to detect the user issuing a request and an *instance identifier* to detect the workflow instance that the request is related to. This is because several instances of the same workflow may be running at the same time and they may have different execution histories (e.g., an instance 1 of the invoicing process where only $t1$ was executed by $u1$ and an instance 2 where both $t1$ was executed by $u1$ and $t2$ was executed by $u2$). *Invariants* indicate values that should not be modified during a workflow instance execution.

Data-flow properties are obtained by using differential analysis, i.e. comparing the differences in the data values between traces, as is done in related work (e.g., [53, 44]). For each trace, the analysis compares the values of all parameters in each request in relation to (i) the same parameter in other requests of the same trace, (ii) the same parameter in other traces, (iii) other parameters in the same trace, and (iv) other parameter in other traces. AEGIS does not apply syntactic annotation (as, e.g., [44]) to identify the data type of each parameter, and does not try to discover possible values or intervals for data elements, because it does not enforce particular values that were seen in the traces (except for *constants*). Below, we describe the differential analysis used to identify each kind of data-flow property.

Let WS be the set of traces τ_i used for analysis, each τ_i be composed of requests r_j and responses s_j , and each request or response have a set P of parameters (k, v) . Considering the same request r_j in every trace $\tau \in WS$, if a parameter (k, v) appears in only a strict subset $\tau' \subset \tau$ of the traces, it is considered optional and ignored, i.e. dropped from the URL in the labeling function L . *Constants* are parameters that are present in every trace $\tau \in WS$ for the same URL of a request r_j and whose key k and value v never change. An example is the parameter `action=create`, which is in $t1$ of traces τ_1, τ_2 , and τ_3 in Example 1. *Don't cares* are parameters that appear in every trace $\tau \in WS$ for the same URL of a request r_j and whose key k remains constant, but whose value v is different in at least one of the requests. One example is `prod=abc`, `prod=def` and `prod=ghi` in $t1$ of Example 1 annotated as `prod=<<DC>>`. An *instance identifier* is a key k whose value v is present in every request r of a trace τ , with different v 's in every trace. In Example 1, the parameter `id` is an instance identifier, since it has the value 1 in every request of τ_1 , the value 2 in every request of τ_2 , and the value 3 in every request of τ_3 . Notice that what must remain constant is the value and not the key, so it is possible to have an instance identifier called, e.g., `id` in one request and `iid` in another request. A *user identifier* is a parameter that comes from a response issued by the server, is stored in a cookie, sent in every request of a trace and whose value changes in every trace in WS . In Example 1, only URLs are shown in the traces, but the cookie `sid` is sent with every request, as can be seen towards the end of the example. *Invariants* are values v that remain constant during a trace, change between traces in WS and are not present in every request of a trace (as opposed to instance identifiers). Two examples are the `value` parameter in $t1$ and $t3$ in Example 1 and the `token` in $t4$ and $t5$ of Example 2. Like instance identifiers, invariant

values should not change, but their *keys* might, so that an invariant can be called, e.g., **price** in one request and **amount** in another. There may be many invariants in a workflow, so they are annotated as $\langle\langle I_1 \rangle\rangle, \langle\langle I_2 \rangle\rangle$, for run-time enforcement (there may be several don't cares too, but they are not enforced and do not need separate annotations).

The result of Step 1 is a tuple (PN, L) , where PN is a Petri net obtained from process mining and L is a labeling function that associates to each transition in the net a URL annotated with the identified data properties. Although the inferred model (PN, L) is obtained automatically, it can be edited by a user before being sent for monitor synthesis. Control-flow constraints can be changed by graphically adding or removing places or transitions in the Petri net (or tasks and gateways in BPMN), while data properties can be modified by adding or removing annotations on the URLs.

3.2 Step 2 - Monitor synthesis

Monitor synthesis takes as input the tuple (PN, L) obtained from Step 1 and, optionally, augments it with security properties given by the user. As an example, the user can specify a set of authorization constraints $\text{SoD}(tx, ty)$ indicating that tasks tx and ty must be executed by different users (the same goes for other constraints, such as BoD). The user must also indicate whether the monitor should enforce an authorization policy, which will be specified at run-time. A symbolic transition system is obtained from the augmented tuple and sent to a model checker, which computes a graph containing all valid executions of the workflow. The graph represents these executions compactly by using symbolic states, symbolic users, and sharing common paths.

Security properties specification. All behaviors of the web application that satisfy the specified security properties (namely those deriving from authorization policies and constraints) are represented by the executions of a symbolic transition system $S = (V, Tr)$, where V is a set of state variables and Tr is a set of transitions. In general, each workflow task corresponds to one transition. The enforced properties, and as a consequence the variables in V , fall into four categories. Therefore, V can be seen as the union of four disjoint sets V_{CF} , V_{DF} , V_C , and V_A , explained below. First, *control-flow* constraints (involving state variables from the set V_{CF}) are automatically derived from the Petri net PN by using Boolean variables p_i 's, one for each place in PN , indicating the presence or absence of tokens¹⁰ in these places; and Boolean variables d_{ti} 's, one for each task, representing the fact that a task has been executed or not. Second, data values of parameters annotated as invariants are represented by variables v_i and g_i (g stands for ghost) in V_{DF} . Data types are abstracted and every v_i and g_i is represented by an integer, since any type can be encoded as an integer. Third, the set V_C contains Boolean functions h_t 's, one for each task, keeping track of which user has executed task t . The functions start with a value False for every transition and user, which is updated after each task execution. *Authorization constraints* of the form $\text{SoD}(tx, ty)$ are represented by an enabling condition $\neg h_{tx}(u)$ in transition ty . BoD constraints can be encoded in a similar way as SoD and cardinality constraints can be specified by using a function $\text{count}(u)$ that keeps track of the number of tasks executed by each user. Fourth, an *authorization policy* is represented by constraints

on Boolean functions a_t 's, one for each task, that involve state variables from the set V_A and return True iff a user is authorized to perform task t . The functions a_t 's are an interface to the authorization policy that is provided at run-time. Although we do not detail them here, other security policies can be encoded in this framework, such as data-based access control and Chinese Wall [55] by using separate authorization and history functions for data objects.

Transitions in Tr have the shape

$$t(u) : en_{CF} \wedge en_C \wedge en_A \rightarrow act_{CF} || act_C || act_A || act_{DF}$$

where $t(u)$ is an identifier, the *en*'s are predicates on the state variables in V representing the *enabling conditions* of the transitions (in terms of control-flow, constraints, and authorization policy, respectively), and the *act*'s are parallel ($||$) assignments to the variables in V representing the *effects* of executing a transition (again, for each security property). Data variables in V_{DF} are not used in the conditions, only in the assignments of transitions that contain data invariants, as $g_i := v_i$. The values of v_i are taken as input at run-time. As an example, the transition for task $t2$ in Example 1 is

$$t2(u) : p1 \wedge \neg d_{t2} \wedge \neg h_{t1}(u) \wedge a_{t2}(u) \rightarrow \\ p1, p2, d_{t2}, h_{t2}(u) := F, T, T, T$$

indicating that, for this transition to be executed, there must be a token in $p1$, $t2$ should not have been executed ($\neg d_{t2}$), the user u should not have executed $t1$ ($\neg h_{t1}(u)$) and the same user u should be authorized to execute $t2$ ($a_{t2}(u)$); the result of its execution is that a token is removed from $p1$, placed in $p2$ and the functions d_{t2} and h_{t2} are updated to record that $t2$ has been executed and user u has executed $t2$, respectively. Since $t2$ does not contain invariants, there is no assignment to data values. However, $t1$ contains $g_1 := v_1$ in the update, where the value of v_1 will be taken as the value of parameter **value** of the incoming request at run-time.

Monitor synthesis. The transition system S is fed to a symbolic model checker, which computes a reachability graph RG representing all possible executions of the workflow by a set of symbolic users. A procedure \mathcal{MS} to compute this graph is described in [6]. \mathcal{MS} is based on backward reachability, starting from a goal formula describing the termination of the workflow and applying transitions until a fix-point is reached. RG is a directed graph whose edges are labeled by task-user pairs in which users are symbolically represented by variables (called *user variables*) and whose nodes are labeled by a symbolic representation (namely, a formula of first-order logic) of the set of states from which it is possible to reach a state in which the workflow successfully terminates. A Datalog [11] program M (with negation) is derived from RG by generating a clause of the form $\text{can_do}(u, t) \leftarrow \beta_n$ for each node n in the graph. An invariant $d_t \Rightarrow v_i = g_i$, for every v_i in the assignments of transition t , is conjoined with each clause β_n . This invariant specifies that after the execution of each transition the value of a variable remains the same as the value of its respective ghost variable. M is then translated to SQL (aggregation-free SQL and non-recursive Datalog with negation are equivalent and the translation is straightforward [46]) and the SQL program is capable of answering—after being instantiated with a concrete authorization policy—user requests to execute tasks in a workflow in such a way that the authorization and execution constraints are not violated, the authorization policy is respected and termination of the workflow is guaranteed, thus enforcing the specified security

¹⁰In Section 3.2, *token* refers to Petri net tokens instead of security tokens

properties and solving the run-time WSP.

The result of the monitor synthesis step is a tuple (M, L) , where M is the monitor generated from RG and L is the labeling function, which now maps from transitions in the system to annotated HTTP requests.

3.3 Step 3 - Run-time monitoring

Step 3 takes as input (M, L) and, if previously specified, an authorization policy TA specifying which users can execute which tasks. The authorization policy is used to populate a database queried by M , resulting in a concrete monitor.

A reverse proxy intercepts all incoming requests to the application and decides, for each request, whether it is part of a workflow or not. To do so, it tries to match the URL and parameters in the request to annotated URLs and parameters stored in L , taking into account the constant, ignored and don't care values. If there is no match, the proxy forwards the request to the application, as it is not part of any workflow. If there is a match, the proxy associates the request to a task t of a workflow $W(T, U)$ and checks the annotated URL for $\langle\langle IID \rangle\rangle$ and $\langle\langle UID \rangle\rangle$ values, extracting the instance i and the user u . The user identifier is a cookie value that must be mapped to a user name in the policy. This is done by capturing login actions, storing the cookies issued to each user, and later retrieving the user names based on the cookie.

To enforce data invariants, when the proxy receives a request for the first URL containing the annotation $inv = \langle\langle I_i \rangle\rangle$, it stores the value of the parameter inv as v_i . When any subsequent task containing $\langle\langle I_i \rangle\rangle$ is accessed, the value of the incoming annotated parameter inc is compared to the stored value ($v_i = inc$). In Example 1, for instance, $t1$ sets the value of parameter $value$, while $t3$ checks that this value is unchanged. In these requests, the monitor query $can_do(u, t) \leftarrow \beta$ is dynamically conjoined with the data invariant condition, becoming $can_do(u, t) \leftarrow \beta \wedge v_i = inc$. The WSP solution remains unchanged because the monitor still guarantees termination if the invariants are respected.

Finally, the proxy issues a request $can_do(u, t)$ to the monitor of instance i of W and acts based on its response by either forwarding the request or dropping it.

4. EVALUATION

AEGIS was implemented in Python 2.7. We capture execution traces as Zest¹¹ scripts exported by ZAP, extract data properties from them, aggregate them into an XES [47] file and use ProM [47] to output a PNML [28] file containing the mined Petri net. The implementation of the monitor synthesis algorithm is the one from [6], using the MCMT model checker [24], which takes as input a PNML file and outputs the synthesized SQL monitor. We use mitmproxy¹² and instantiate it with the generated monitor as in the example below:

```
mitmdump -R http://localhost:80 -p8080 -s httpmonitor.py
```

where `mitmdump -R` starts the proxy in the reverse mode to accept requests on port 8080, process them using the `httpmonitor.py` script and, possibly, forward them to the application (`http://localhost:80`). The `httpmonitor.py` script intercepts requests and responses using a proxy API, performs the URL matching, queries a MySQL database

(where the authorization policies are stored) by using the synthesized queries, and either forwards or drops the request. The proxy also supports HTTPS connections.

4.1 Experimental setup

We tested AEGIS on popular open-source applications (shown in Table 1), synthesizing monitors in the configurations $\langle\mathcal{P}, \mathcal{C}, \mathcal{I}\rangle$ and $\langle\mathcal{I}\rangle$. Applications 1-4 are ERP platforms, 5-6 are e-health applications and 7-10 are e-commerce applications. Column *Application* contains the name of each application; *Language* shows the language in which it was developed; *Params* describes the predominant method used for parameter passing (although an application can use several methods) and *Downloads* reports the number of downloads (applications 1-6) or public installations (applications 7-10).

The different languages show the versatility of the black-box approach, which has to be tailored to support each parameter passing method (to annotate and match URLs). Supporting new applications that use the same method is straightforward, whereas supporting new methods (e.g., OData [40]) requires new functionality for model inference. The number of downloads and installations is a measure of the popularity of the applications and it comes from official repositories (applications 2, 3, 5, and 6), data in the web page of the project (applications 1 and 4), or related work [39] (applications 7-10). The number of actual deployments for applications 1-6 is not available as they are usually internal to an organization and not indexed by search engines. The numbers shown for applications 7-10 were obtained using Google dorks and are from 2014 [39].

We pre-configured the applications using demo data (e.g., financial accounts for ERP, products for e-commerce, patients for e-health) either available during installation or generated by us. We then captured four execution traces for each workflow (as in [53]) and two login traces for each application.

To compare AEGIS in different ERP applications, we used workflows offered by all of them: *Purchase order* (PO), *Sales order* (SO), *Purchase invoice* (PI), and *Sales invoice* (SI). They are slightly different in each application, varying from 4 to 6 tasks, usually with a gateway defining 2 to 3 alternative execution branches. Figure 4 shows at the top the patient visit workflow mined from OpenEMR (where the lines labeled by = represent BoD constraints added by us). The same Figure shows at the bottom the lab analysis workflow mined from BikaLIMS. In these 6 applications, we added the authorization constraints and specified policies with 10 users assigned to each task, generating $\langle\mathcal{P}, \mathcal{C}, \mathcal{I}\rangle$ monitors. The number of users was arbitrarily chosen because it influences the time taken to answer queries (discussed in Section 4.2).

Table 1: Applications used in the experiments

#	Application	Language	Params	Downloads
1	Odoo	Python	JSON	2M
2	Dolibarr	PHP	GET	850k
3	WebERP	PHP	GET	617k
4	ERPNext	Python	JSON	25k
5	OpenEMR	PHP	GET	382k
6	BikaLIMS	Python	REST	111k
7	OpenCart 1.5.3.1	PHP	GET	9M
8	TomatoCart 1.1.7	PHP	GET	119k
9	osCommerce 2.3.1	PHP	GET	80k
10	AbanteCart 1.0.4	PHP	GET	21k

¹¹<https://goo.gl/jNyFK4>

¹²<https://mitmproxy.org/>

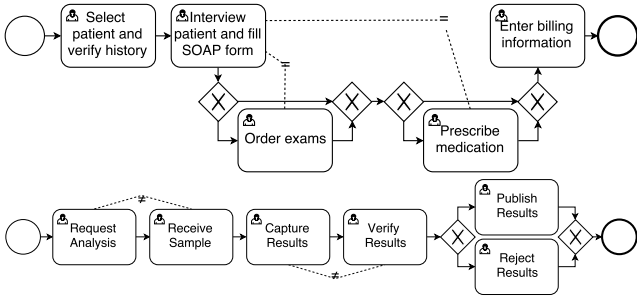


Figure 4: Patient visit workflow from Open-EMR (top) and lab analysis workflow mined from BikaLIMS (bottom)

The workflows for e-commerce applications are similar to the one shown in Figure 3. For these applications, we use the $\langle I \rangle$ configuration, thus neither constraints nor authorization policies were defined. Applications 7 and 8 have a vulnerability allowing attackers to shop for free due to improper validation of PayPal Express Checkout tokens, which can be replayed from previous transactions, as explained in Section 2.2 (CVE-2012-4934 for TomatoCart). Applications 9 and 10 allow an attacker to buy products and pay to himself, by tampering with a parameter that indicates who should receive the payment for a PayPal Payments Standard transaction (CVE-2012-2991 for osCommerce).

All applications were deployed as Docker [34] containers and the tests as Selenium scripts, using the architecture described in [13], which allows us to achieve repeatable experiments by automatically testing the applications in five steps: (i) start a new container with the application; (ii) run the workflow in the Selenium script without monitoring; (iii) start the monitor; (iv) run the workflow with monitoring; (v) capture results and destroy the container. The experiments ran on a MacBook 2014 laptop with a 1.3GHz dual-core Intel Core i5 processor and 8GB of RAM.

4.2 Results

The enforcement of security properties and mitigation of vulnerabilities was successful in all applications, which was confirmed by manual inspection. In applications 1-6, we tested the enforcement of policies and constraints by trying the attacks described in Section 3 (workflow bypass, workflow violations, and authorization violations). The monitor was able to block situations such as the same user executing an entire workflow (SoD violation), and users trying to access tasks that were not assigned to them. In applications 7-10, we tried to exploit the vulnerabilities described above. In applications 7-8, the attacks were unsuccessful because `token` was detected as an invariant and automatically enforced. In applications 9-10, the `PayeeId` parameter was detected as a constant, since every trace in the input was related to the same shop. Constant values are usually not enforced, only used to match URLs (details in Section 3). For applications 9-10, we edited the inferred model by annotating `PayeeId` with `<<I>>`, so that requests with any value of `PayeeId` are controlled by the monitor, and used invariant enforcement with a constant, instead of with the first received value, to check that in every request containing `PayeeId`, its value is equal to `ShopId` (the constant obtained in the traces). Manual editing could be avoided by doing inference from a

Table 2: Monitoring overhead

#	App.	Synth.	Original	Query	Aegis	Overhead
1	Odoo					
	PO	21.3 s	112 ms	6 ms	132 ms	20 ms
	SO	22.4 s	170 ms	7 ms	213 ms	43 ms
	PI	14.3 s	174 ms	7 ms	213 ms	39 ms
	SI	17.9 s	104 ms	7 ms	116 ms	12 ms
2	Dolibarr					
	PO	14.2 s	93 ms	5 ms	103 ms	10 ms
	SO	14.3 s	92 ms	4 ms	104 ms	12 ms
	PI	13.2 s	89 ms	5 ms	97 ms	8 ms
	SI	14.7 s	90 ms	5 ms	105 ms	15 ms
3	WebERP					
	PO	20 s	51 ms	6 ms	59 ms	8 ms
	SO	21.1 s	50 ms	5 ms	57 ms	7 ms
	PI	18.3 s	30 ms	6 ms	37 ms	7 ms
	SI	19.5 s	32 ms	4 ms	39 ms	7 ms
4	ERPNext					
	PO	13.3 s	222 ms	7 ms	251 ms	29 ms
	SO	12.9 s	327 ms	14 ms	411 ms	84 ms
	PI	15.9 s	263 ms	10 ms	327 ms	64 ms
	SI	13.7 s	272 ms	13 ms	318 ms	46 ms
5	OE	19.1 s	95 ms	7 ms	112 ms	17 ms
6	BL	31.2 s	306 ms	7 ms	326 ms	20 ms
7	OpC	19.1 s	65 ms	6 ms	77 ms	12 ms
8	TC	15.8 s	63 ms	4 ms	71 ms	8 ms
9	osC	22.2 s	79 ms	7 ms	95 ms	16 ms
10	AC	19.8 s	117 ms	8 ms	127 ms	10 ms

dataset containing execution traces of different shops.

We measured the overhead of the monitors in terms of model inference and monitor synthesis and by comparing the execution of each workflow with and without monitoring. Each execution was repeated 10 times and Table 2 shows the results. Column *Appl.* shows the application under test (and the specific workflow tested for ERP applications); *Synth.* shows the median time to infer a model from the captured traces and synthesize a monitor for each workflow. *Original* reports the median time between receiving a request and sending a response with no monitor (measured by mitmproxy without the `httpmonitor.py` script); *Query* reports the median time for the monitor to answer to a query (ignoring the time taken by the proxy to invoke the script, translate an incoming request to a monitor query, forward the request, etc); *Aegis* reports the median time of a response with the monitor script (the time taken by the application, plus the translation time, plus the querying time); and *Overhead* shows the overhead incurred by the use of the monitor as seen by a user (the difference between *Aegis* and *Original*).

The time in column *Query* varies with the size of a workflow and the number of users and constraints, as reported in [6], which describes a linear growth due to the *LogSpace* complexity of the queries used. The time in column *Aegis* adds, to the time in *Query*, the time to process and match URLs, which depends on the data structures used.

As shown in Table 2, the overhead varied from 8ms to 84ms, with a median of 13.5ms, out of which less than 10ms in most cases is spent in querying the monitor. The overhead variability is due to the complexity of the workflows and the time taken to translate a request to a monitor. For instance, applications 1 and 4 have a large overhead because of the time to flatten JSON requests. Monitor synthesis is computationally much more expensive, but it is run only once for each workflow (unless there is a change in the application).

A modular approach allows the scalability of the synthesis procedure, which was tested with workflows of up to 500 tasks and returned a monitor in less than 10 minutes [16]. We did not test the performance of monitoring concurrent executions of workflows. Since there is no interaction between instances, we believe that any additional overhead would be related to request processing in the proxy and database access to answer monitor queries.

5. DISCUSSION AND LIMITATIONS

Model inference. There are solutions to design web applications as workflows, as well as frameworks that allow their declarative description [48]. However, model-driven development of web applications is not common. This highlights the need for model inference, which can be static or dynamic [10, 25, 26, 51]. The use of dynamic analysis by process mining allows us to develop a black-box approach, but we intend to investigate a hybrid solution, combining static and dynamic techniques, enriched with the knowledge of the user.

Authorization. AEGIS provides an easy way to enforce authorization policies and constraints on the actions a user can trigger when interacting with a service (via URL requests). Obtaining the same behavior within an application is not trivial. In fact, it must be done differently for each application and the granularity of the permissions therein offered may not be easily related to URL requests. As an example, the granularity of permissions in the applications we experimented with varied from actions in a module (e.g., creating an invoice in the finance module of Dolibarr) to binary module access (i.e. users that can access a module can perform all of its actions). Though the former provides permissions over actions, maybe not all actions in the process are covered, e.g., *Send by email* of Figure 2. In the latter, it would not be possible to create SoD constraints between actions within the same module. None of the applications we tested supports authorization based on individual URLs nor authorization constraints. Moreover, the policy enforced by AEGIS can be applied on top of the existing one (if any) and can be easily specified by connecting tasks (obtained from the HTTP traces) to users (obtained from the application).

Other approaches. Some properties enforced by AEGIS can be achieved with other tools, e.g., Web Application Firewalls (WAFs) [4], which filter incoming requests based on user-defined rules, or Run-time Application Self-Protection [?]. Although the use of WAFs is a best practice, they are typically incapable of handling logical flaws [15]. Indeed, most tools protect against injection vulnerabilities rather than business logic flaws [14] and a defense-in-depth approach, combining multiple tools, is recommended [1]. We are unaware of any single tool that encompasses all the properties enforced by AEGIS (related run-time enforcement approaches are discussed in the next section). Security mechanisms can be coded directly in the application, but this is error-prone, unscalable, and difficult for legacy applications.

Usability. We have not tested AEGIS with end-users, but we believe the learning curve is not steep, since most steps are automated. The manual steps—trace capturing, (optional) model editing, and (optional) authorization policy definition—are facilitated by the use of well-known tools for test automation and business process editing.

Limitations. AEGIS only sees the traffic between users and the target application, ignoring messages between a third party and a user, and between the application and a

third party. Model inference ignores some message exchange formats, e.g., XML. The first limitation is architectural; the second is an implementation issue. It is possible to have a parser for each format that returns (k, v) pairs and searches for them in incoming requests. The invariants that AEGIS detects and enforces are only exact matches, which is the most common case in web applications [53]. Supporting more complex relations between data values (e.g., values in a list) requires new analysis and new data annotations.

We did not see false positives (i.e. blocking valid requests) during our evaluation, because we inferred models from a diverse dataset and tested them to ensure they captured all the executions we foresaw. However, AEGIS is not immune to false positives caused by poorly inferred models that do not match all executions of an application. Therefore, the user may not want AEGIS to block incoming requests, which could prevent legal executions. Instead, it can be used for soft enforcement, where denied requests represent deviations that must be logged so that a human agent can later examine them.

AEGIS synthesizes monitors that work in isolation, disregarding any possible inter-workflow and inter-instance dependencies and constraints. Related works consider such constraints when executing applications in several tabs [8].

6. CONCLUSION

We have described, implemented, and evaluated AEGIS, a technique and tool to enforce authorization policies and constraints, control- and data-flow integrity and ensure the satisfiability of web applications. We have tested our implementation with relevant open-source applications. The experiments clearly show the validity of our approach in enforcing the desired properties and mitigating related vulnerabilities. The performance results show an acceptable overhead incurred at run-time.

Related work. There are many works related to the enforcement of authorization, control- and data-flow integrity (separately) in web applications, and mitigation of related vulnerabilities, such as missing authorization checks [43, 54]; workflow and state violations [3, 12]; logic vulnerabilities [19]; forceful browsing [8]; and parameter tampering [7, 2]. These approaches are white-box, whereas AEGIS is completely black-box. A black-box approach to block request forgery was described in [31, 30], but it ignores data-flow and authorization. Black-box enforcement of control- and data-flow integrity has been done in Ghostrail [9] by dynamically replicating on the server-side valid user clicks, form entries, links, and parameters. Ghostrail does not consider authorization, and needs a fresh replica for each user session, which is not scalable.

Process mining has been used to optimize user interaction [41], but we are unaware of any work using it to enforce security in web applications. Web application workflow models have been used to detect anomalous user behavior [33] and to find logic vulnerabilities by capturing execution traces, identifying behavioral patterns, and generating test cases [39]. AEGIS is focused on enforcement, so these techniques are complementary. It is possible to, e.g., find a vulnerability using [39] and mitigate it using AEGIS (this is what was done with the e-commerce vulnerabilities). The enforcement of authorization constraints for collaborative web applications was studied in [20, 22, 21]. The authors considered different application domains than us, there was no discussion about workflow satisfiability, and their evaluation was limited to

prototypical applications.

The closest related works are BLOCK [32] and InteGuard [53]. Both use a reverse proxy, construct control- and data-flow policies using invariants detected from network traces, and rely on manual identification of critical requests. BLOCK also extracts invariants from session information in PHP applications. InteGuard is tailored for multi-party application integration, where most steps are automatic (not human tasks) and workflows must be executed from beginning to end in one shot. Neither tool enforces authorization policies nor constraints. FlowWatcher [37] uses a similar approach of external monitoring, but enforces only authorization policies specified in a domain specific language. **Future work.** We intend to test AEGIS in more real-world applications, e.g., those in [53, 44], and to extend it with code analysis to measure the coverage of inferred models by executing traces and checking followed paths. We would also like to explore monitor inlining [23], which requires source code changes to embed the monitor into the application.

7. REFERENCES

- [1] N. Antunes and M. Vieira. Defending against web application vulnerabilities. *Computer*, 45(2):66–72, Feb 2012.
- [2] M. Balduzzi, C.T. Gimenez, D. Balzarotti, and E. Kirda. Automated discovery of parameter pollution vulnerabilities in web applications. In *Proc. of NDSS*, 2011.
- [3] D. Balzarotti, M. Cova, V. Felmetsger, and G. Vigna. Multi-module vulnerability analysis of web-based applications. In *Proc. of CCS*, 2007.
- [4] M. Becher. *Web Application Firewalls*. VDM Verlag, Saarbrücken, Germany, 2007.
- [5] E. Bertino, E. Ferrari, and V. Atluri. The specification and enforcement of authorization constraints in workflow management systems. *TISSEC*, 2(1):65–104, 1999.
- [6] C. Bertolissi, D. R. dos Santos, and S. Ranise. Automated synthesis of run-time monitors to enforce authorization policies in business processes. In *Proc. of ASIACCS*, 2015.
- [7] P. Bisht, T. Hinrichs, N. Skrupsky, R. Bobrowicz, and V.N. Venkatakrishnan. Notamper: automatic blackbox detection of parameter tampering opportunities in web applications. In *Proc. of CCS*, 2010.
- [8] B. Braun, P. Gemein, H.P. Reiser, and J. Posegga. Control-flow integrity in web applications. In *Proc. of ESSoS*, 2013.
- [9] B. Braun, C. Gries, B. Petschkuhn, and J. Posegga. Ghostrail: Ad hoc control-flow integrity for web applications. In *Proc. of IFIP SEC*, 2014.
- [10] R.P. Castillo, I.G.R. de Guzman, and M. Piattini. Business process archeology using marble. *Inf. and Soft. Tech.*, 53(10):1023 – 1044, 2011.
- [11] S. Ceri, G. Gottlob, and L. Tanca. What You Always Wanted to Know About Datalog (And Never Dared to Ask). *TKDE*, 1(1):146–166, 1989.
- [12] M. Cova, D. Balzarotti, V. Felmetsger, and G. Vigna. Swaddler: An approach for the anomaly-based detection of state violations in web applications. In *Proc. of RAID*, 2007.
- [13] S. Dashevskiy, D. R. dos Santos, F. Massacci, and A. Sabetta. Testrex: a testbed for repeatable exploits. In *Proc. of CSET*, 2014.
- [14] G. Deepa and P.S. Thilagam. Securing web applications from injection and logic vulnerabilities: Approaches and challenges. *Inf. and Soft. Tech.*, 74:160–180, 2016.
- [15] M. Dermann, M. Dziadzka, B. Hemkemeier, A. Hoffmann, A. Meisel, M. Rohr, and T. Schreiber. OWASP Best Practices: Use of Web Application Firewalls, 2008. Available at: https://www.owasp.org/index.php/Best_Practices:_Web_Application_Firewalls.
- [16] D. R. dos Santos, S. Ranise, and S. E. Ponta. Modular synthesis of enforcement mechanisms for the workflow satisfiability problem: Scalability and reusability. In *Proc. of SACMAT*, 2016.
- [17] A. Doupé, L. Cavedon, C. Kruegel, and G. Vigna. Enemy of the state: A state-aware black-box web vulnerability scanner. In *Proc. of USENIX Sec.*, 2012.
- [18] F. Duchene, S. Rawat, J. Richier, and R. Groz. Ligre: Reverse-engineering of control and data flow models for black-box xss detection. In *Proc. of WCRE*, 2013.
- [19] V. Felmetsger, L. Cavedon, C. Kruegel, and G. Vigna. Toward automated detection of logic vulnerabilities in web applications. In *Proc. of USENIX Sec.*, 2010.
- [20] P. Gaubatz, W. Hummer, U. Zdun, and M. Strembeck. Supporting customized views for enforcing access control constraints in real-time collaborative web applications. In *Proc. of ICWE*, 2013.
- [21] P. Gaubatz, W. Hummer, U. Zdun, and M. Strembeck. Enforcing entailment constraints in offline editing scenarios for real-time collaborative web documents. In *Proc. of SAC*, 2014.
- [22] P. Gaubatz and U. Zdun. Supporting entailment constraints in the context of collaborative web applications. In *Proc. of SAC*, 2013.
- [23] G. Gheorghe and B. Crispo. A survey of runtime policy enforcement techniques and implementations. Technical report, University of Trento, 2011.
- [24] S. Ghilardi and S. Ranise. Mcmt: A model checker modulo theories. In *Proc. of IJCAR*, 2010.
- [25] W.G.J. Halfond. Identifying inter-component control-flow in web applications. In *Proc. of ICWE*, 2015.
- [26] W.G.J. Halfond, S. Anand, and A. Orso. Precise interface identification to improve testing and analysis of web applications. In *Proc. of ISSSTA*, 2009.
- [27] S. Hallé, T. Ettema, C. Bunch, and T. Bultan. Eliminating navigation errors in web applications via model checking and runtime enforcement of navigation state machines. In *Proc. of ASE*, 2010.
- [28] L.M. Hillah, E. Kindler, F. Kordon, L. Petrucci, and N. Treves. A primer on the petri net markup language and iso/iec 15909-2. In *Proc. of CPN*, 2009.
- [29] J. Holderer, R. Accorsi, and G. Müller. When four-eyes become too much: a survey on the interplay of authorization constraints and workflow resilience. In *Proc. of SAC*, 2015.
- [30] K. Jayaraman, G. Lewandowski, P.G. Talaga, and S.J. Chapin. Enforcing request integrity in web applications. In *Proc. of DBSec*, 2010.

- [31] K. Jayaraman, P.G. Talaga, G. Lewandowski, S.J. Chapin, and M. Hafiz. Modeling user interactions for (fun and) profit: preventing request forgery attacks on web applications. In *Proc. of PLOP*, 2009.
- [32] X. Li and Y. Xue. Block: a black-box approach for detection of state violation attacks towards web applications. In *Proc. of ACSAC*, 2011.
- [33] X. Li, Y. Xue, and B. Malin. Detecting anomalous user behaviors in workflow-driven web applications. In *Proc. of SRDS*, 2012.
- [34] K. Matthias and S.P. Kane. *Docker: Up & Running*. O'Reilly, 2015.
- [35] M. Meucci and A. Muller. *Testing Guide 4.0*. The OWASP Foundation, 2015.
- [36] T. Murata. Petri nets: properties, analysis and applications. *Proc. of the IEEE*, 77(4):541–580, 1989.
- [37] D. Muthukumaran, D. O’Keeffe, C. Priebe, D. Eysers, B. Shand, and P. Pietzuch. Flowwatcher: Defending against data disclosure vulnerabilities in web applications. In *Proc. of CCS*, 2015.
- [38] OWASP. *OWASP Top 10 - 2013: The Ten Most Critical Web Application Security Risks*. The OWASP Foundation, 2013.
- [39] G. Pellegrino and D. Balzarotti. Toward black-box detection of logic flaws in web applications. In *Proc. of NDSS*, 2014.
- [40] M. Pizzo, R. Handl, and M. Zurmuehl. Odata version 4.0 part 1: Protocol. Technical report, OASIS, 2014. Available at: https://www.oasis-open.org/committees/tc_home.php?wg_abbrev=odata.
- [41] N. Poggi, V. Muthusamy, D. Carrera, and Rania Khalaf. Business process mining from e-commerce web logs. In *Proc. of BPM*, 2013.
- [42] M. Schur, A. Roth, and A. Zeller. Mining workflow models from web applications. *TSE*, 41(12):1184–1201, 2015.
- [43] S. Son, K.S. McKinley, and V. Shmatikov. Fix me up: Repairing access-control bugs in web applications. In *Proc. of NDSS*, 2013.
- [44] A. Sudhodanan, A. Armando, L. Compagna, and R. Carbone. Attack patterns for black-box security testing of multi-party web applications. In *Proc. of NDSS*, 2016.
- [45] F. Sun, L. Xu, and Z. Su. Static detection of access control vulnerabilities in web applications. In *Proc. of USENIX Sec.*, 2011.
- [46] G. Terracina, N. Leone, V. Lio, and C. Panetta. Experimenting with recursive queries in database and logic programming systems. *Theory Pract. Log. Program.*, 8(2):129–165, 2008.
- [47] W.M.P. van der Aalst. *Process Mining*. Springer, 2011.
- [48] E. Vervaet. *The Definitive Guide to Spring Web Flow*. Apress, Berkely, CA, USA, 2008.
- [49] Q. Wang and N. Li. Satisfiability and resiliency in workflow authorization systems. *TISSEC*, 13(4):40:1–40:35, 2010.
- [50] R. Wang, S. Chen, and X. Wang. Signing me onto your accounts through facebook and google: a traffic-guided security study of commercially deployed single-sign-on web services. In *Proc. of IEEE S&P*, 2012.
- [51] W. Wang, Y. Lei, S. Sampath, R. Kacker, R. Kuhn, and J. Lawrence. A combinatorial approach to building navigation graphs for dynamic web applications. In *Proc. of ICSM*, 2009.
- [52] M. Weske. *Business Process Management*. Springer, 2007.
- [53] L. Xing, Y. Chen, X. Wang, and S. Chen. Integuard: Toward automatic protection of third-party web service integrations. In *Proc. of NDSS*, 2013.
- [54] F. Yamaguchi, C. Wressnegger, H. Gascon, and K. Rieck. Chucky: exposing missing checks in source code for vulnerability discovery. In *Proc. of CCS*, 2013.
- [55] F. Yan and P.W.L. Fong. Efficient irm enforcement of history-based access control policies. In *Proc. of ASIACCS*, 2009.